

Binary Search Trees (BST)

- A BST is a *binary tree* that (if not empty) also follows the following storage rule regarding its nodes' items:

- ◆ For any node *n* of the tree, every item in *n*'s *LST*, if not empty, is *less than* the item in *n*
- ◆ For any node *n* of the tree, every item in *n*'s *RST*, if not empty, is *greater than* the item in *n*

NOTE: any node implies rule must hold not just for the overall tree but also for each and every subtree

NOTE: Enforcement of the rule requires that the nodes' items can be compared with the usual comparison operators *<, >, == etc.* (i.e., can be arranged in a single line, from small to large)

Caveat Storage rule for defining BST can differ (*points valid under a certain definition may be invalid under other definitions*):

- LST items → *less than or equal*, RST items → *greater than*
- LST items → *less than*, RST items → *greater than or equal*
- LST items → *less than or equal*, RST items → *greater than or equal*

this is used by textbook but we'll use the above instead

1

Binary Search Trees (BST)

Searching

- Due to the ordering of nodes (based on the values of items at the nodes), searching for a target value in a BST can be greatly faster than in a simple binary tree with no ordering
 - ◆ We know that all items in the LST are less than the root's item and all items in the RST are greater than the root's item
 - ◆ We compare the target value with that of root's item and make a *better informed decision* based on the result of the comparison
 - Such decision making works only if the specified storage rule is enforced
→ without the rule, a match could be found anywhere in the tree

2

Binary Search Trees E.g.: IntSet using a BST

We can implement the IntSet ADT we saw using a BST

```
class IntSet
{
public:
    struct btNode
    {
        int data;
        btNode* left;
        btNode* right;
    };
    ...
    bool contains(int anInt) const;
    bool add(int anInt);
    bool remove(int anInt);
    ...
private:
    btNode* root_ptr;
    ...
};
```

Items are stored in
a tree; tree is to be
maintained as a BST

3

Binary Search Trees E.g.: IntSet using a BST Querying Containment

- The public member function **contains** determines whether a BST contains a target item **anInt**
 - ◆ **bool contains(int anInt) const;**
- Because the items in a BST are ordered according to the rule listed on the first slide, **contains** doesn't have to inspect every item in the BST to ascertain containment
 - ◆ Let **cursor** be a local pointer (initialized to the *root pointer*) used to keep track of the current search position in the BST
 - ◆ **cursor** will move through the BST in search of **anInt**
 - ◆ **cursor** will use the rule to *always move along the path where anInt might occur* (the next slide elaborates this)

4

Binary Search Trees *E.g.: IntSet using a BST* Querying Containment

- There are *four* possibilities at each search position:
 - ◆ **cursor** becomes *0*, indicating that we've moved off the bottom of the tree
 - Run out of nodes where a match can be found
 - **return false;**
 - ◆ **anInt** is *smaller* than the item at **cursor** node
 - **anInt** can only occur in the *left* subtree of the BST rooted at **cursor** node
 - **cursor = cursor->left;**
 - ◆ **anInt** is *larger* than the item at **cursor** node
 - **anInt** can only occur in the *right* subtree of the BST rooted at **cursor** node
 - **cursor = cursor->right;**
 - ◆ **anInt** is *equal* to the item at **cursor** node
 - Match of **anInt** found
 - **return true;**
- It is straightforward to implement **contains** with a loop

5

Binary Search Trees *E.g.: IntSet using a BST* Adding an Item

- The public member function **add** inserts an entry **anInt** (received as a parameter) into a BST
 - ◆ **bool add(int anInt);**
- The **add** function first deals with the special case in which the tree is *empty*
 - ◆ **root_ptr = new btNode;**
root_ptr->data = anInt;
root_ptr->left = root_ptr->right = 0;
- If the tree is *not empty*, the function searches the tree to see if a spot for adding **anInt** as a leaf can be found
 - ◆ Again, let local pointer **cursor** (initialized to the root pointer) be used to keep track of the current search position in the BST

6

Binary Search Trees *E.g.: IntSet using a BST* Adding an Item

- If **anInt** is *equal to* the data at **cursor** node
 - ◆ Entry already exists
 - **return false;**

7

Binary Search Trees *E.g.: IntSet using a BST* Adding an Item

- If **anInt** is *less than* the data at **cursor** node
 - ◆ Check the **left** field at **cursor** node
 - If it's **0**, create a new node containing **anInt** and make the **left** field of **cursor** node point to it (and task is done):

```
cursor->left = new btNode;  
cursor->left->data = anInt;  
cursor->left->left = cursor->left->right = 0;
```
 - If it's *not 0*, move the cursor to the **left** and continue the search for a correct spot to insert the entry:

```
cursor = cursor->left;
```

8

Binary Search Trees *E.g.: IntSet using a BST* Adding an Item

■ If **anInt** is *greater than* the data at **cursor** node

◆ Check the **right** field at **cursor** node


➤ If it's **0**, create a new node containing **anInt** and make the **right** field of **cursor** node point to it (and task is done):

```
cursor->right = new btNode;  
cursor->right->data = anInt;  
cursor->right->left = cursor->right->right = 0;
```

➤ If it's *not 0*, move the cursor to the *right* and continue the search for a correct spot to insert the entry:


```
cursor = cursor->right;
```

Binary Search Trees *Inserting a Node*



Binary Search Trees


Inserting a Node



11

Binary Search Trees


Inserting a Node



12

Binary Search Trees

Inserting a Node



13

Binary Search Trees

E.g.: `IntSet` using a BST
Removing an Item

- The public member function **remove** removes a specified item **anInt** (if exists) from a BST
 - ◆ `bool remove(int anInt);`
- To implement **remove** *directly* entails *many special cases* to be dealt with and a *precursor* (similar to that used when removing an item from a singly linked list) be maintained
 - ◆ We will (next) consider a popular *indirect* way
 - (often referred to as the “usual standard way” or something similar)
 - ◆ It uses two helper functions (next)

14

Binary Search Trees

Removing a Node Helper Functions

(Each receives, among others, root pointer of BST as *reference parameter* **bst_root**)

■ **bst_remove** BST may be *empty* or *non-empty*

- ◆ If the target **anInt** was in the BST, then **anInt** has been removed, **bst_root** now points to the root of the new (smaller) BST, and the function returns true
- ◆ If the target **anInt** was not in the BST, then the BST is unchanged, and the function returns false

■ **bst_remove_max** For *non-empty* BST only

- ◆ The *largest item* in the BST has been removed, and **bst_root** now points to the root of the new (smaller) BST
- ◆ A copy of item removed is returned via a *reference parameter*

15

Binary Search Trees

Removing a Node **bst_remove** Helper

■ **bst_remove** has a recursive implementation to remove the target **anInt**

- ◆ *Tree* could be *empty* → function simply returns false
- ◆ Target **anInt** could be *less than* the *root* → make a *recursive call to the left*
 - `bst_remove(bst_root->left, anInt);`
- ◆ Target **anInt** could be *greater than* the *root* → make a *recursive call to the right*
 - `bst_remove(bst_root->right, anInt);`
- ◆ Target **anInt** could be *equal* to the *root* → ... (next)

16

Binary Search Trees

Removing a Node **bst_remove** Helper

Target Equal to Root

- If root node's value is equal to target **anInt**, we have a match
- To remove this node, we have to consider two situations:
 - ◆ *Root has no children or only 1 child*
 - (both children are empty or only 1 child is non-empty)
 - ◆ *Root has two children*
 - (both children are non-empty)

17

Binary Search Trees

Removing a Node **bst_remove** Helper

Target Equal to Root

Root Has No Children or Only 1 Child

- In this case we can delete the root node and make 0 (if no children) or the root of the non-empty child (if only 1 child) the new root node
- This requires three steps:
(1st and 3rd steps are identical for the different cases)

	<i>Has No Children</i>	<i>Has Only Left Child</i>	<i>Has Only Right Child</i>
❶	<code>old_bst_root = bst_root;</code>		
❷	<code>bst_root = 0;</code>	<code>bst_root = bst_root->left;</code>	<code>bst_root = bst_root->right;</code>
❸	<code>delete old_bst_root;</code>		

18

Binary Search Trees

Removing a Node **bst_remove** Helper

Target Equal to Root and Root Has 2 Children

- In this case we have to delete the root node and make an *appropriate child* the new root node
- We'll *replace root w/ the node w/ the largest value in LST*
 - ◆ Call **bst_remove_max** on the *LST*
 - Have it remove the node with the largest value from the LST
 - Have it set **bst_root->data** (passed as *reference parameter*) equal to the largest value (that was in the node removed from the LST)
 - ◆ **bst_remove_max(bst_root->left, bst_root->data);**

We could just as well
replace root w/ the node w/ the smallest value in RST
by calling a corresponding **bst_remove_min** on the *RST*

19

Binary Search Trees

Removing a Node **bst_remove_max** Helper

Outline of (Recursive) Algorithm

- If (tree has *no right child*) // largest item @ root (base case)
 - ◆ Copy root node's data into the **data**-field reference parameter
 - ◆ Delete root node and make left child (may be 0) the new root
 - (using 3 steps similar to that given for **bst_remove**)
- Else // tree has *right child* (largest item not @ root)
 - ◆ Make recursive call to delete largest item from RST
 - Passing along the **data**-field *reference parameter*
 - (side effect to be delivered via the parameter is fulfilled in base case above)

20

Binary Search Trees Balanced vs Unbalanced

- There is no requirement that a BST be full, complete, balanced, *etc.*
 - ◆ Only that the rule listed on the first slide be met
- That means it is possible to have a BST that's very unbalanced (for instance, only a single node in the LST but 1000 nodes in the RST)

21


Binary Search Trees Balanced vs Unbalanced

- If the tree (and the subtrees) has drastically more nodes to one side over the other, the performance improvement in searching is diminished
 - ◆ In the worst case, it's no better than a linked list
- It'd be ideal if the node counts are equal (or within 1 node) for each and every corresponding subtree pair, . . .
 - (often described as *weight-balanced*, more on this to come)
 - . . . so that each decision roughly cuts in half the number of nodes that remain to be searched

22

Binary Search Trees Balanced vs Unbalanced


7 comparisons to find the number 3



23

Binary Search Trees Balanced vs Unbalanced

3 comparisons to find the number 3



24

Binary Search Trees Balanced vs Unbalanced

- A *weight-balanced* BST is one in which each subtree has either the same number of nodes as the other or is off by only 1 node
 - ◆ This property holds for *any node* in the tree
- This type of tree offers the best performance improvement for searches
 - ◆ Each decision "prunes" as many nodes as possible
- However, it is typically too expensive to maintain a weight-balanced BST
 - ◆ (motivates *height-balanced* BSTs such as AVL trees)

25

Textbook Readings

- Chapter 10
 - ◆ Section 10.5

26